

MongoDB GridFS

GridFS is a framework to store & access large set of data. It divides the data into chunks and store them into different documents.

- API Provided by MongoDB for storing large files such as audio, video and images.

- Package that can be plucked into any application to make storing large files easier

Provides a way for storing large files in database instead of in the file system.



MongoDB GridFS



Problem: In MongoDB document size is limited to 16 MB.



Gridfs Solves the size limitation problem

1. Breaks the files to smaller manageable chunks
2. Stores these chunks of data in one collection called **fs.chunks**
3. Stores the information about the whole file itself in another collection called **fs.files**
4. Connects these documents by properties that are references to each other



MongoDB GridFS

fs.chunks collection

1. The size of each chunk is 255KB
2. No. of chunks created depends on the file size
3. Chunks stores the actual data.
4. Each chunk is linked to the file information by “files_id” property.
5. The “files_id” points to a document that is stored in fs.files collection.



```
MongoDB Enterprise > db.fs.chunks.find({}, {data:0}).pretty();
{
  "_id" : ObjectId("6093e3130ee4d7e7115c5ea2"),
  "files_id" : ObjectId("6093e3120ee4d7e7115c5ea1"),
  "n" : 0
}
{
  "_id" : ObjectId("6093e70fe8a48c9a7a484db0"),
  "files_id" : ObjectId("6093e70fe8a48c9a7a484daf"),
  "n" : 0
}
{
  "_id" : ObjectId("6093e70fe8a48c9a7a484db1"),
  "files_id" : ObjectId("6093e70fe8a48c9a7a484daf"),
  "n" : 1
}
{
  "_id" : ObjectId("6093e70fe8a48c9a7a484db2"),
  "files_id" : ObjectId("6093e70fe8a48c9a7a484daf"),
  "n" : 2
}
```

MongoDB GridFS

fs.files collection contains the information about the file

1. File name
2. Average size of each chunk
3. Upload date
4. Size of file (in bytes)
5. File metadata

```
MongoDB Enterprise > db.fs.files.find().pretty();
{
  "_id" : ObjectId("6093e3120ee4d7e7115c5ea1"),
  "length" : NumberLong(130797),
  "chunkSize" : 261120,
  "uploadDate" : ISODate("2021-05-06T12:37:39.852Z"),
  "filename" : "boy.jpg",
  "metadata" : {

  }
}
{
  "_id" : ObjectId("6093e70fe8a48c9a7a484daf"),
  "length" : NumberLong(1679701),
  "chunkSize" : 261120,
  "uploadDate" : ISODate("2021-05-06T12:54:39.149Z"),
  "filename" : "me.jpg",
  "metadata" : {

  }
}
```



MongoDB GridFS

The **mongofiles** utility makes it possible to manipulate files stored in your MongoDB instance in GridFS objects from the command line.

The mongofiles tool is part of the MongoDB Database Tools package.

<https://www.mongodb.com/try/download/database-tools>

Run mongofiles from the system command line, not the mongo shell.



MongoDB GridFS

mongofiles <options> <connection-string> <command> <filename or _id>

Options. You may use one or more of these options to control the behaviour of mongofiles.

Connection String. The connection string of the mongod to connect to with mongofiles.

Command - Use one of these commands to determine the action of mongofiles.

Filename - name of the file to be saved in the data base

C:\....\bin>mongofiles put me.jpg --db=img

C:\....\bin>mongofiles --help

C:\....\bin>mongofiles get me.jpg --db=img



files.find().pretty();

MongoDB GridFS

Once the file is stored in the database fs.files & fs.chunks collection can be used to get the information about the file.

>db.fs.files.find().pretty();

```
MongoDB Enterprise > db.fs.files.find().pretty();
{
  "_id" : ObjectId("6093e3120ee4d7e7115c5ea1"),
  "length" : NumberLong(130797),
  "chunkSize" : 261120,
  "uploadDate" : ISODate("2021-05-06T12:37:39.852Z"),
  "filename" : "boy.jpg",
  "metadata" : {
  }
}
```

>db.fs.chunks.find().pretty();

```
i87q0+Ko16/Jqbz38dvHo0v8AX6r-fp/4j2ptfjHDj8+mH8Twzr4U6CmPx6JfFp
enSuGmyW0a9Da9Vr6Lc2/wCI9lc0rH+o/n0/c6s14/8AFdKvC+TVF9tfR4+ba
G2hfH+jQmrVr16tHp/V+Le/d9e7rT+HpGjjX/L/AMXx6f0t4V8enTpN9F/9jo
69Po0+PVza99er/affu2vdx/ydVTRQ+Fwfmf9nrkL6PVf8ARzb9V+f03/PvWNW0
/ACbp/wB6/wB99PbvbTpjNP1/sdf/2Q=="
}
MongoDB Enterprise > db.fs.chunks.find({}, {data:0}).pretty();
{
  "_id" : ObjectId("6093e3130ee4d7e7115c5ea2"),
  "files_id" : ObjectId("6093e3120ee4d7e7115c5ea1"),
  "n" : 0
}
```



MongoDB GridFS

Getting the no of chunks created for each file:

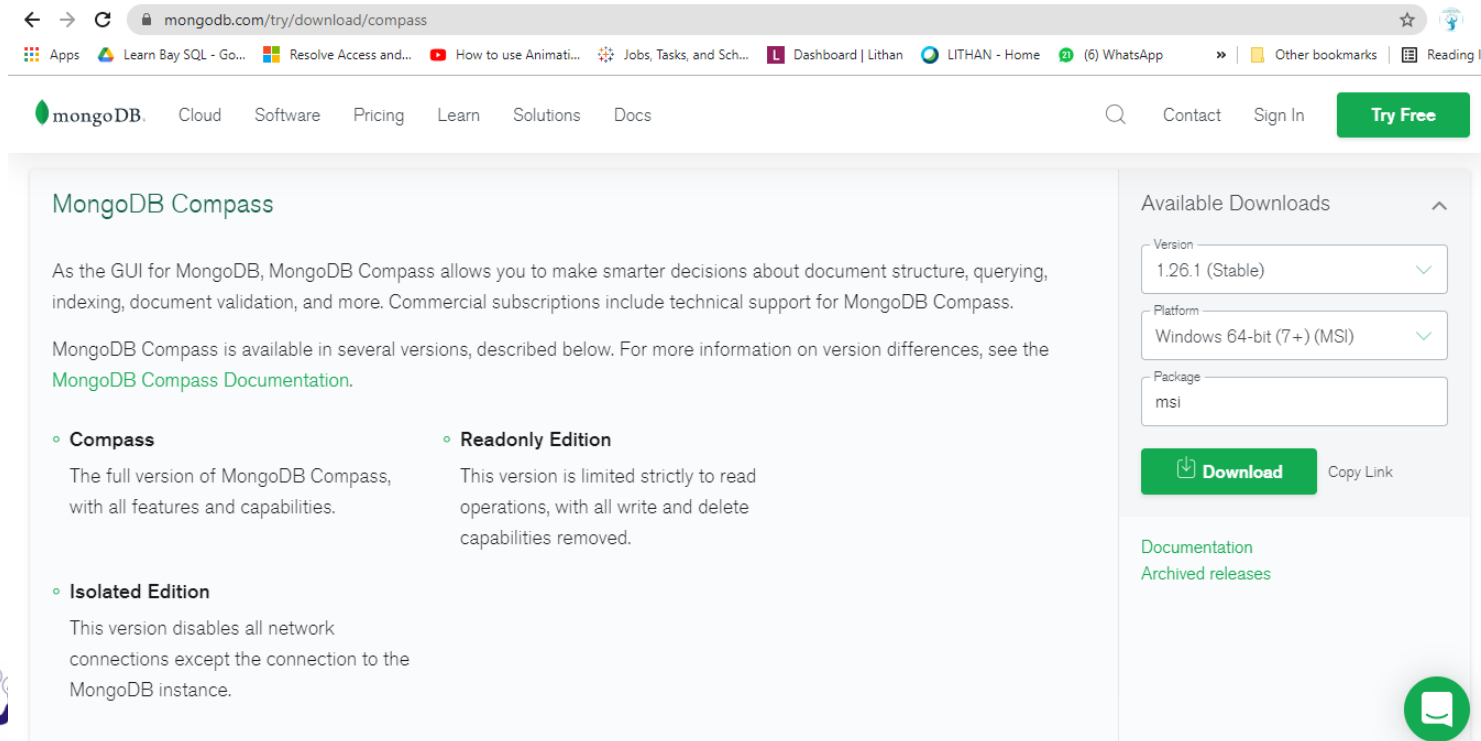
```
>db.fs.chunks.find(  
    {"files_id" : ObjectId("64a7932d07bbfa87bef08aa6")}  
).count()
```



MongoDB Compass

It's a GUI Interface for handling MongoDB database. It is a convenient tool for performing all CRUD operations without manually writing queries. It helps in many activities such as indexing, document validation, etc.

<https://www.mongodb.com/try/download/compass>



The screenshot shows the MongoDB Compass download page. The browser address bar displays 'mongodb.com/try/download/compass'. The page header includes the MongoDB logo and navigation links: Cloud, Software, Pricing, Learn, Solutions, Docs, Contact, Sign In, and a green 'Try Free' button. The main content area is titled 'MongoDB Compass' and describes it as the GUI for MongoDB. It lists three versions: Compass, Readonly Edition, and Isolated Edition, each with a brief description of its capabilities. On the right side, there is a section for 'Available Downloads' with dropdown menus for Version (1.26.1 (Stable)), Platform (Windows 64-bit (7+) (MSI)), and Package (msi). Below these is a green 'Download' button and a 'Copy Link' option. At the bottom right, there are links for 'Documentation' and 'Archived releases'.

MongoDB Compass

As the GUI for MongoDB, MongoDB Compass allows you to make smarter decisions about document structure, querying, indexing, document validation, and more. Commercial subscriptions include technical support for MongoDB Compass.

MongoDB Compass is available in several versions, described below. For more information on version differences, see the [MongoDB Compass Documentation](#).

- **Compass**
The full version of MongoDB Compass, with all features and capabilities.
- **Readonly Edition**
This version is limited strictly to read operations, with all write and delete capabilities removed.
- **Isolated Edition**
This version disables all network connections except the connection to the MongoDB instance.

Available Downloads

Version: 1.26.1 (Stable) ✓

Platform: Windows 64-bit (7+) (MSI) ✓

Package: msi

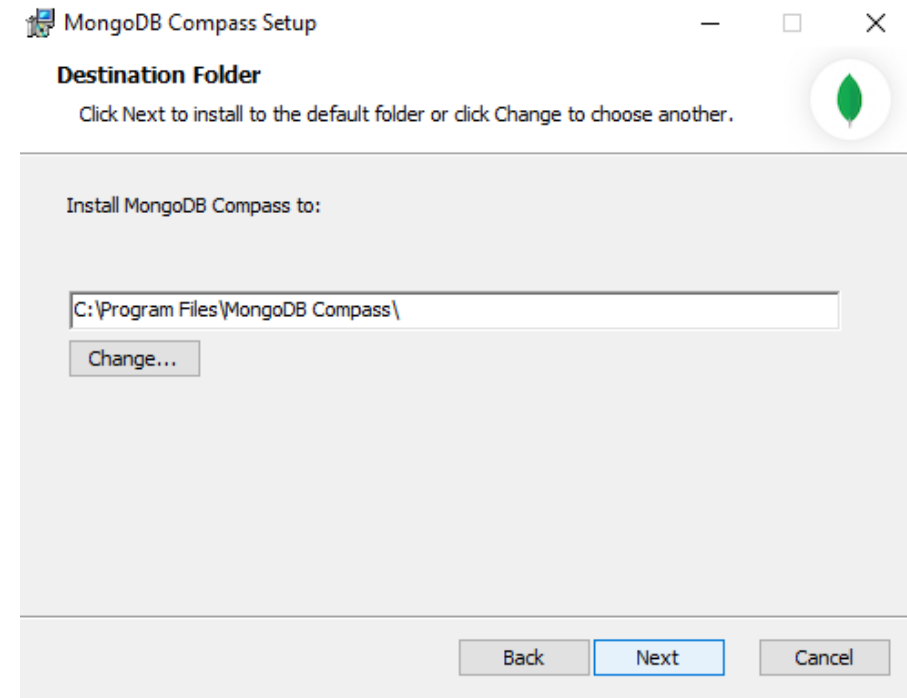
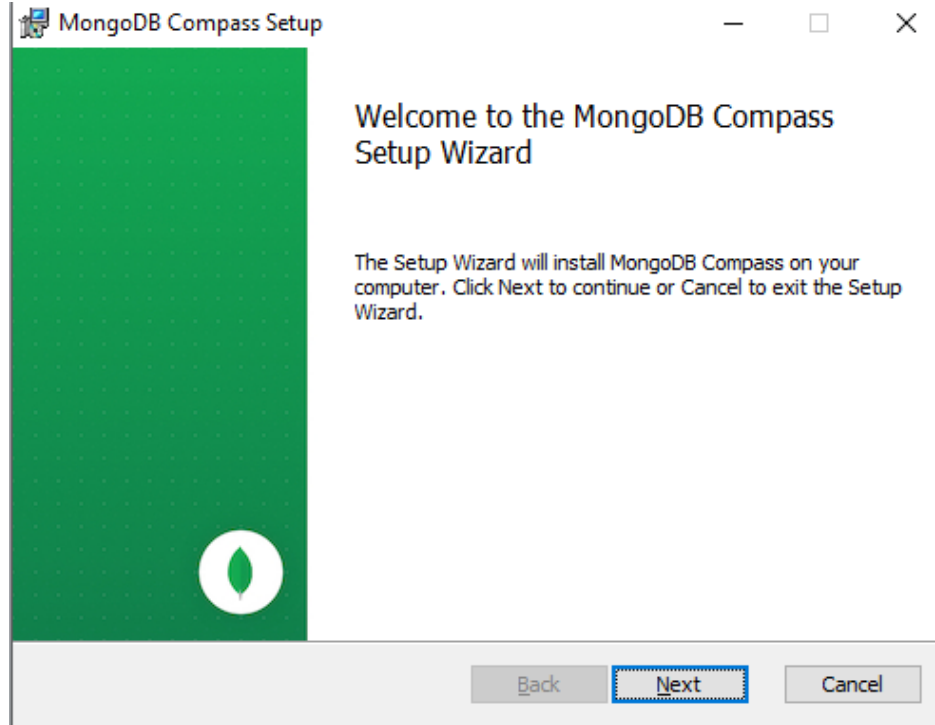
[Download](#) [Copy Link](#)

[Documentation](#)
[Archived releases](#)




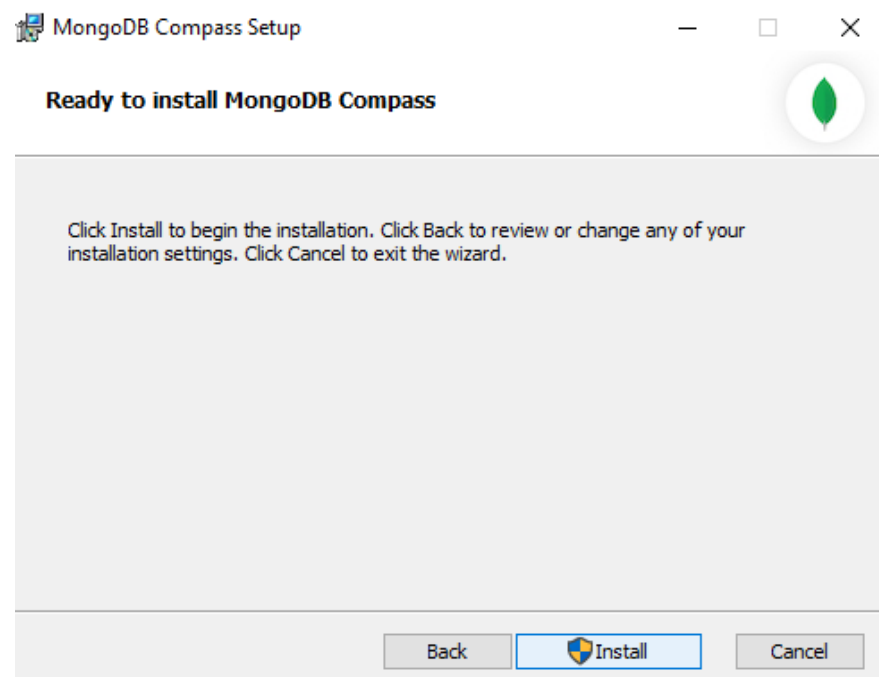
MongoDB Compass

Double click on the downloaded file to initiate the installation:



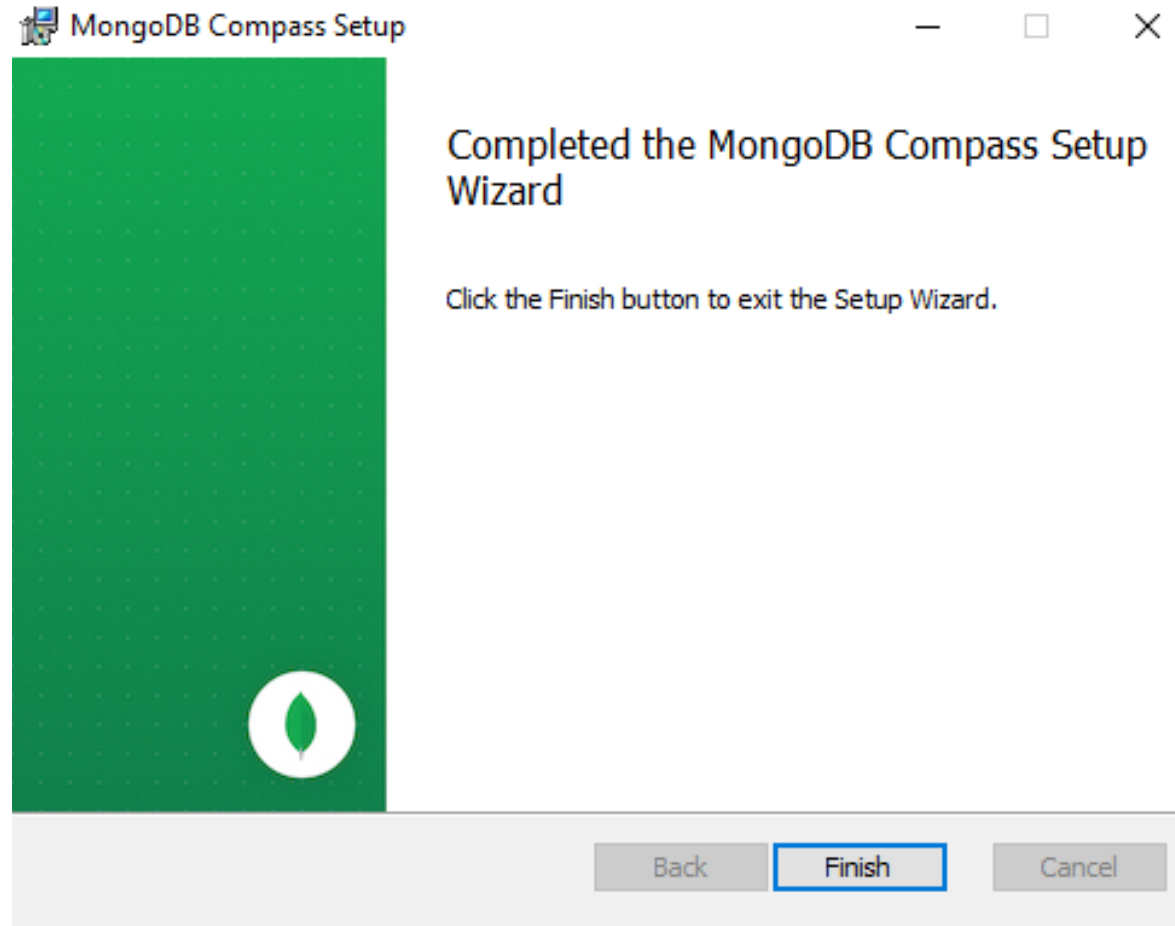
MongoDB Compass

Click on Install proceed with the installation:



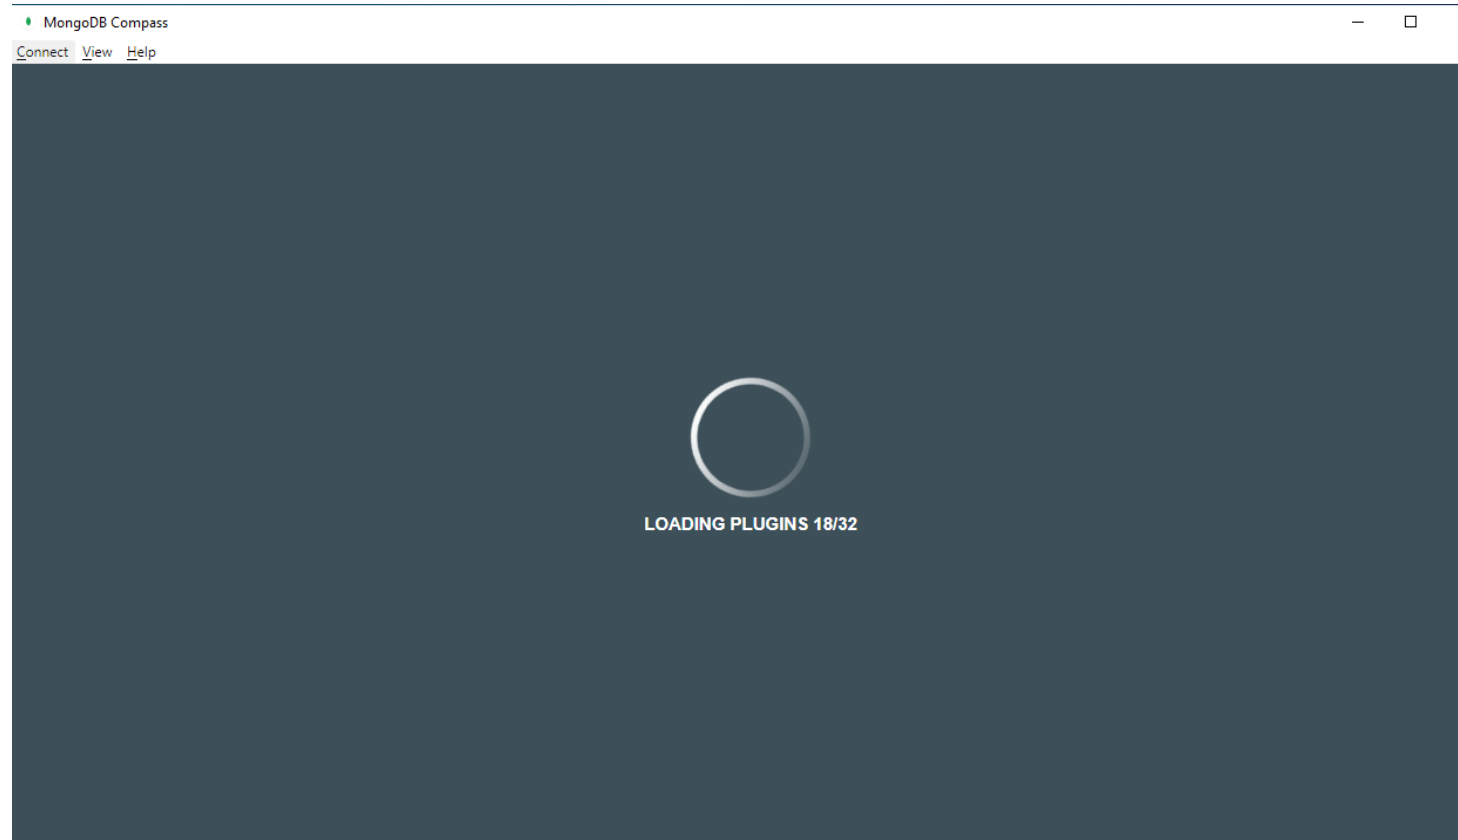
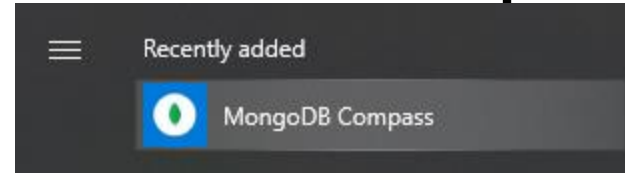
MongoDB Compass

Click on Finish to complete the installation:



MongoDB Compass

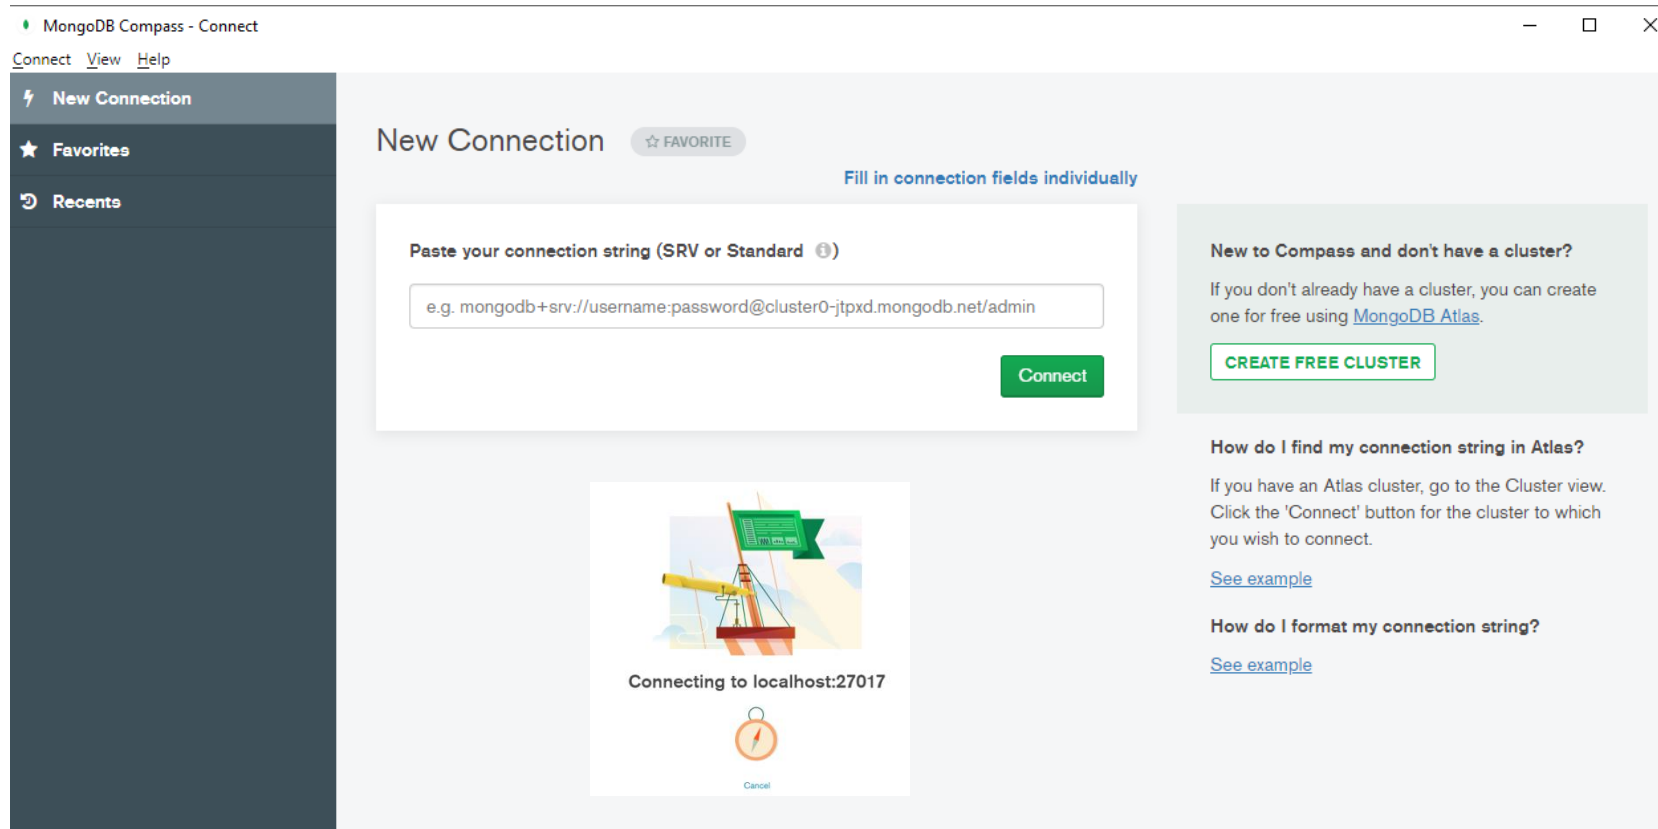
Go to Start Menu and click on



MongoDB Compass

Connect to the cluster (server). We will be connecting to the local host.

Provide the connection string and click on connect



NOTE: Mongo server must be running on your system to connect to the host



MongoDB Compass

Once connected you will be able to see all the databases in the mongo

MongoDB Compass - localhost:27017

Connect View Help

Local

9 DBS 40 COLLECTIONS

☆ FAVORITE

HOST
localhost:27017

CLUSTER
Standalone

EDITION
MongoDB 4.4.5 Enterprise

Filter your data

- > admin
- > batch1
- > batch2
- > config
- > demo
- > img
- > local
- > mydb
- > test

+

Databases Performance

CREATE DATABASE

Database Name ^	Storage Size	Collections	Indexes	
admin	32.0KB	0	1	
batch1	224.0KB	7	8	
batch2	64.0KB	2	2	
config	12.0KB	0	2	
demo	32.0KB	1	1	
img	5.3MB	2	4	
local	36.0KB	1	1	
mydb	36.0KB	2	2	
test	2.7MB	25	30	



MongoDB Compass

MongoDb Compass allows us to perform below operations:

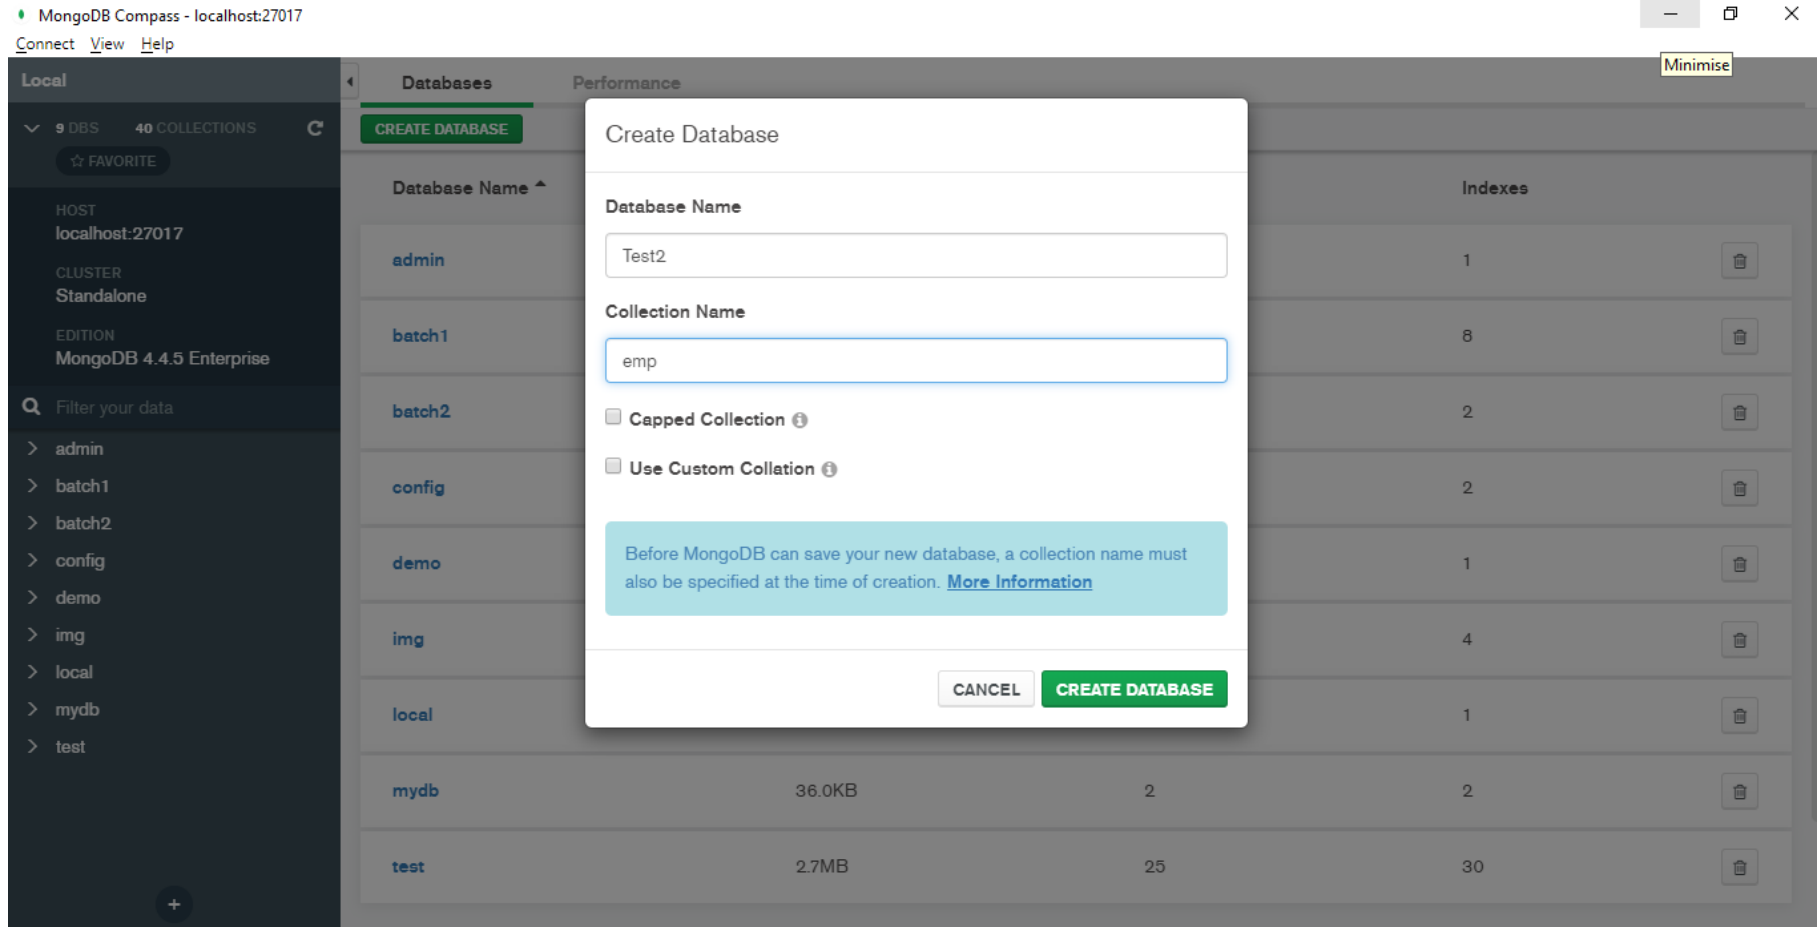
- Create Database
- Create Collection
- Perform CRUD operations
- Add data to collection
 - Import File
 - Insert Document
- Use Options
 - Filter
 - Project
 - Sort
- Aggregations
- Indexes



dation

MongoDB Compass

Create Database



MongoDB Compass

Add data to collection

MongoDB Compass - localhost:27017/test.salary

Connect View Collection Help

Local

9 DBS 40 COLLECTIONS

☆ FAVORITE

HOST
localhost:27017

CLUSTER
Standalone

EDITION
MongoDB 4.4.5 Enterprise

Filter your data

- > admin
- > batch1
- > batch2
- > config
- > demo
- > img
- > local
- > mydb
- test
 - abc
 - abcd

test.salary Documents

Documents

FILTER {EID : 1

ADD DATA

Insert to Collection test.salary

VIEW {} ≡

```
1 //**
2 * Paste one or more documents here
3 */
4 {
5   "_id": {
6     "$oid": "6097ca9d6fed09228188b68f"
7   }
8 }
```

CANCEL INSERT

56 TOTAL SIZE 4.3KB AVG. SIZE 78B INDEXES 1 TOTAL SIZE 32.0KB AVG. SIZE 32.0KB

Validation

OPTIONS FIND RESET ↺ ...

Displaying documents 1 - 1 of 1 < > ↻ REFRESH



MongoDB Compass

Aggregations

MongoDB Compass - localhost:27017/test.salary

Connect View Collection Help

Local

9 DBS 40 COLLECTIONS

☆ FAVORITE

HOST
localhost:27017

CLUSTER
Standalone

EDITION
MongoDB 4.4.5 Enterprise

Filter your data

- > admin
- > batch1
- > batch2
- > config
- > demo
- > img
- > local
- > mydb
- ▼ test
 - abc
 - abcd

test.salary

Aggregations

DOCUMENTS 55 TOTAL SIZE 4.2KB AVG. SIZE 78B INDEXES 1 TOTAL SIZE 32.0KB AVG. SIZE 32.0KB

Documents Aggregations Schema Explain Plan Indexes Validation

COLLATION Untitled - Modified SAVE SAMPLE MODE AUTO PREVIEW

55 Documents in the Collection

Preview of Documents in the Collection

Select an operator to construct expressions used in the aggregation pipeline stages. [Learn more](#)

Output after \$group stage (Sample of 7 documents)

```
1 /**
2  * _id: The id of the group.
3  * fieldN: The first field name.
4  */
5 {
6   _id: "$DEPT",
7   TOTALCOST: { $sum: "$SALARY" },
8   AvgSal: { $avg: "$SALARY" },
9   Minsal: { $min: "$SALARY" },
10  Maxsal: { $max: "$SALARY" }
11 }
```

Example 1:

_id: ObjectId("608a7474a538f278ddd80e5a")
EID: 1004
DEPT: "MIS"
DESI: "Manager"
SALARY: 134789

Example 2:

_id: ObjectId("608a7474a538f278ddd80e5b")
EID: 1001
DEPT: "OPS"
DESI: "Director"
SALARY: 380000

Example 3:

_id: "HR"
TOTALCOST: 1350000
AvgSal: 150000
Minsal: 150000
Maxsal: 150000

Example 4:

_id: "mis"
TOTALCOST: 153767
AvgSal: 153767
Minsal: 153767
Maxsal: 153767



MongoDB Compass

Validations

MongoDB Compass - localhost:27017/test.salary

Connect View Collection Help

Local

9 DBS 40 COLLECTIONS

☆ FAVORITE

HOST
localhost:27017

CLUSTER
Standalone

EDITION
MongoDB 4.4.5 Enterprise

Filter your data

- > admin
- > batch1
- > batch2
- > config
- > demo
- > img
- > local
- > mydb
- ▼ test
 - abc
 - abod

test.salary Validation

test.salary

DOCUMENTS 55 TOTAL SIZE 4.2KB AVG. SIZE 78B INDEXES 1 TOTAL SIZE 32.0KB AVG. SIZE 32.0KB

Documents Aggregations Schema Explain Plan Indexes Validation

```
1 {  
2   SALARY: 1  
3   $gt: 20000  
4 }  
5 }
```

✓ Sample Document That Passed Validation

```
_id: ObjectId("608a7474a538f278ddd80e5a")  
EID: 1004  
DEPT: "MIS"  
DESI: "Manager"  
SALARY: 134789
```


✗ Sample Document That Failed Validation

No Preview Documents



Sharding MongoDB

Sharding is the process of partitioning your data across multiple servers. It is a type of database partitioning that separates very large database into faster, smaller and more easily manageable parts called shards.



Sharding MongoDB

MongoDB uses the shard key to distribute the collection's documents across shards. The shard key consists of a field or multiple fields in the documents.

Why Sharding ?

Scalable - data is growing continuously

High Availability

Ability to control data distribution

Application Transparent

Cost effective

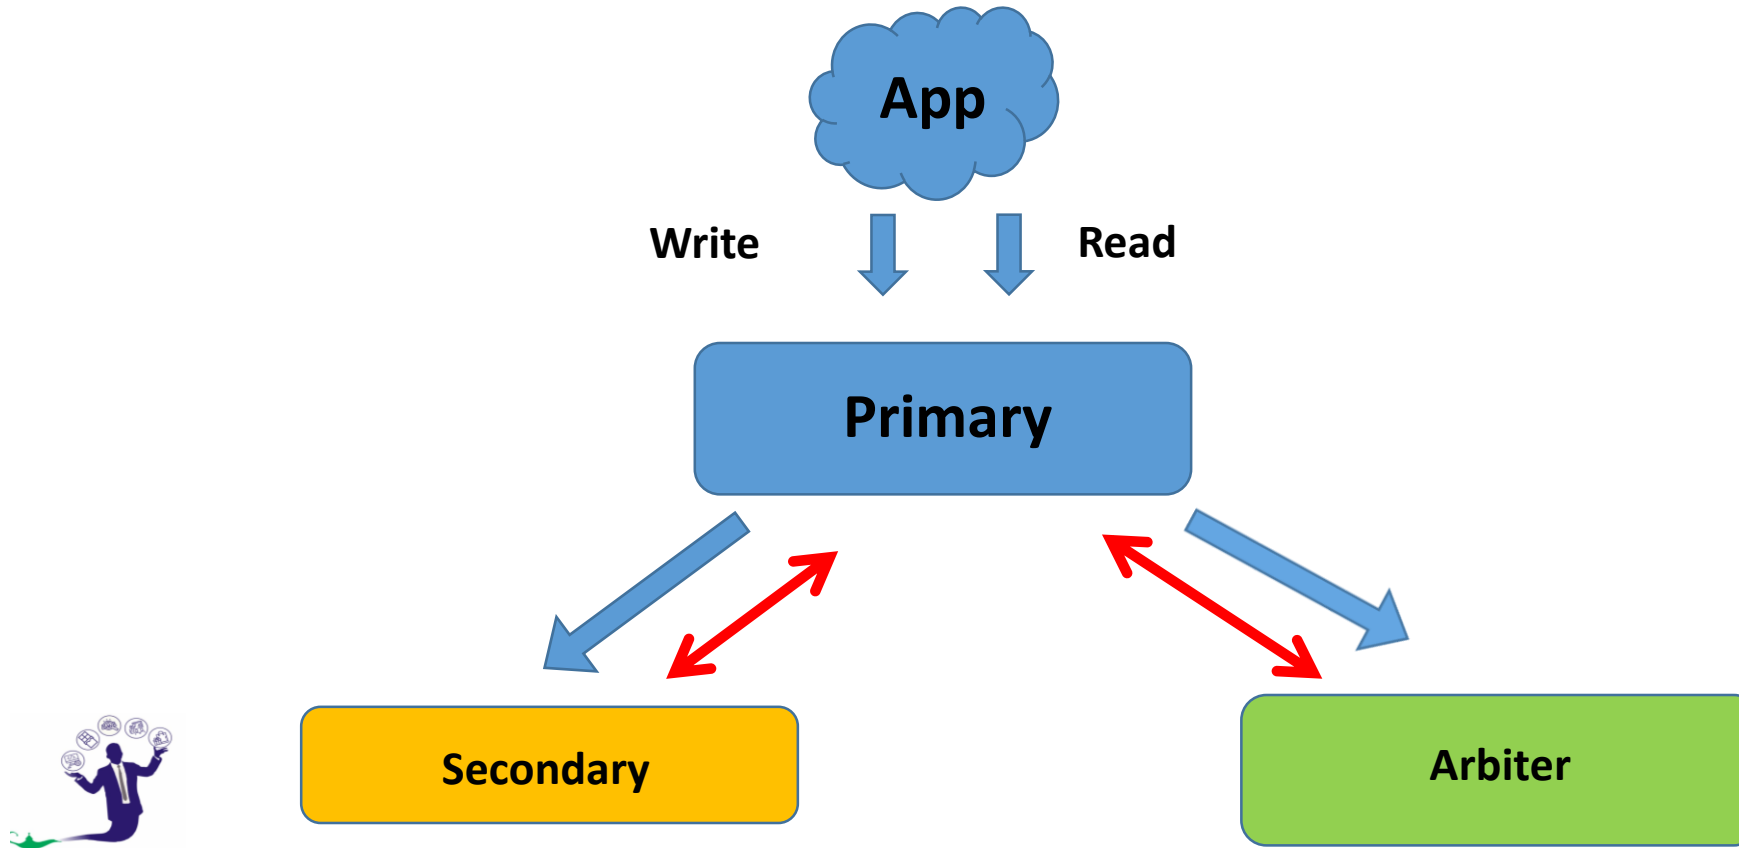
No database downtime



Replication in MongoDB

A replica set is a group of mongod instances (server) that maintain the same data set. A replica set contains several data bearing nodes and optionally one arbiter node. Of the data bearing nodes, one and only one member is deemed the primary node, while the other nodes are deemed secondary nodes.

The primary node receives all write operations



Replication in MongoDB

Major features of replica:


Asynchronous Replication -Secondary replicate the primary's and apply the operations to their data sets asynchronously.

Automatic Failover (electionTimeoutMillis period (10 seconds by default))

Read Preference - can specify a read preference to send read operations to secondaries.

Mirrored Reads – operations can be in the cache of secondary.





Thanks!

***EVERY ENDING
IS REALLY JUST A
NEW BEGINNING***



Rajeev Garg
Data Analytics Trainer